

GNUmed

Antoine Janssens¹, Maria Karegla¹, Michiel Creemers², and Yosra Ben Taarit²

¹Advanced Master's in Cybersecurity

²Master of Artificial Intelligence

October 19, 2025

Abstract

Keywords

Big Data, Health, Privacy, GNU, open-source

1 Introduction

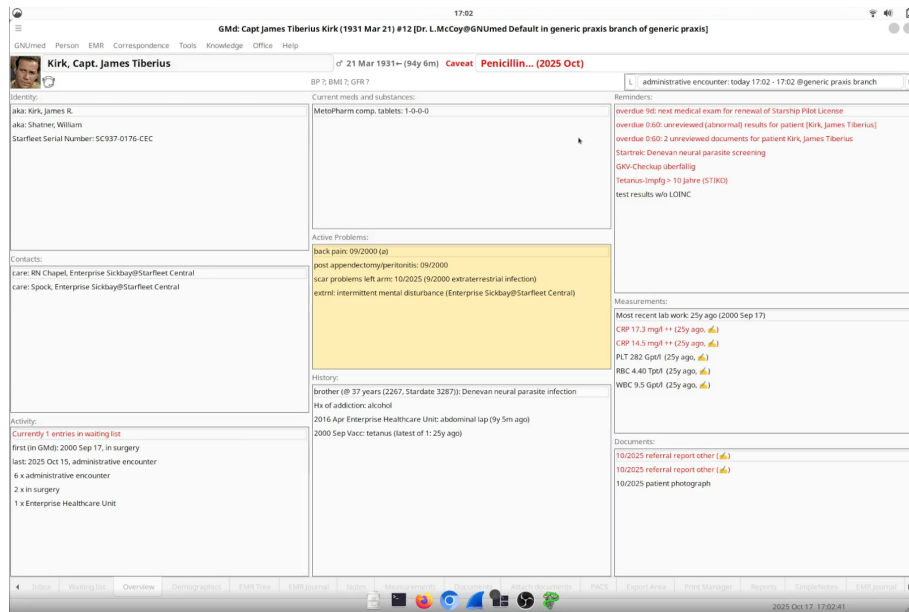


Figure 1: GNUmed Screenshot of patient file.

2 Technical description of the system

2.1 System Architecture

Based on the GNUmed introduction document[?], we deduct that GNUmed makes use of a flexible client-server model. Thigh high-level component diagram details their key responsibilities.

The **GNUmed Client** consists of the main GNUmed application. A Database Service Broker is used to abstract the backend service distribution. This provides a unified data source that the

rest of the application can make use of, providing transparency to the physical layout for user and developer alike. In adherence with security requirements, a client side cryptography engine handles the encryption of sensitive data, ensuring that even before being transmitted, the data of this nature cannot be accessed.

The **GNUmed Backend** is built on a PostgreSQL database and is not monolithic. It is made up of a collection of distributed services. Data that is related is regarded as a service and hence can be treated as a logical component. Services are not necessarily present on the same physical server. Within the database, business logic is enforced through the use of triggers. These triggers also maintain an Audit Trail, as required legally, by automatically capturing all changes, by keeping copies of the unchanged instances. The tables are highly normalized in order to optimize storage. For the client, data access is simplified through an abstraction layer of updateable views, which sit on top of the tables. Access to these data is managed through a user hierarchy that regulates permissions.

Finally, the **Communication** between the client and server is secured using industry standard protocols, with examples being SSL for encryption and Kerberos for authentication. The architecture also allows for **Optional Middleware**, to be placed between client and backend, allowing for more complex integrations or business logic external to the database.

2.2 The Logical Data Model

The system's UML Class Diagram as shown in Figure ??, represents the logical data model and details the structure of the data located in the PostgreSQL backend. The figure is outdated and is lacking information. It provides a blueprint for some of the services, as previously described, with each prefix corresponding to a specific schema and hence a service. In this case we have the Demographic Service, represented by dem., Clinical Service by clin., Document Storage Service represented by blobs., and German Specific Schema represented by de_de.

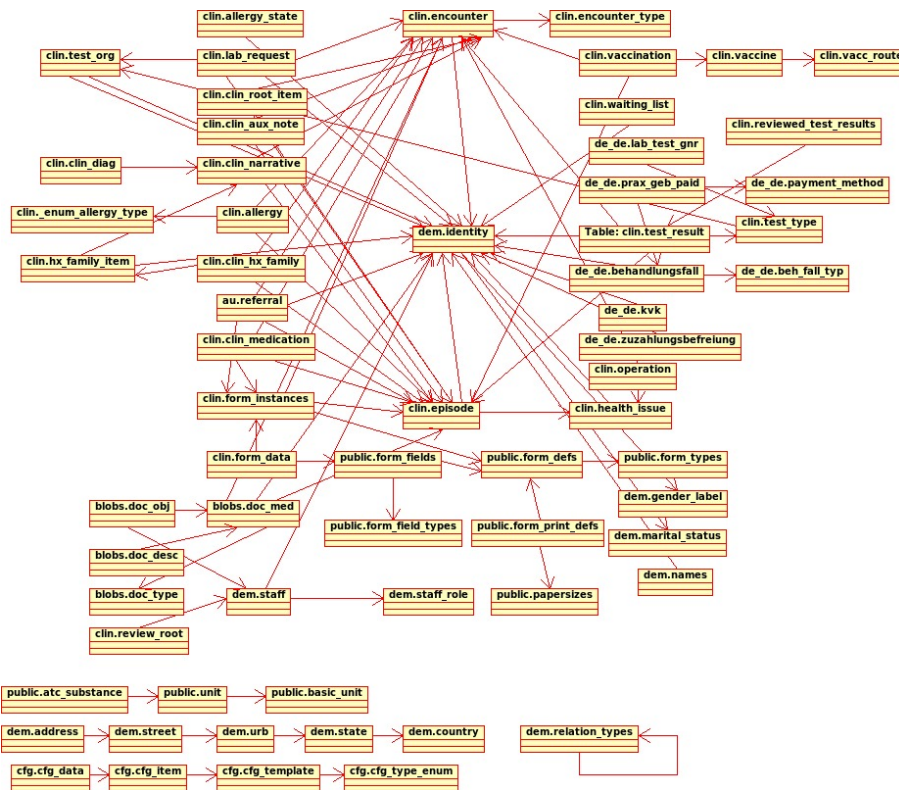


Figure 2: GNUmed UML Class Diagram. Note: outdated (2002) and missing elements [?]

3 technical privacy issues

4 technological recommendations

References

- [1] Res Medicinae. domain.png (927×798).
- [2] GNUmed Development Team. GNUmed-introduction_v4.pdf.